

Data Structures – Week #6

Special Trees

Outline

- Adelson-Velskii-Landis (AVL) Trees
- Splay Trees
- B-Trees

AVL Trees

Motivation for AVL Trees

- Accessing a node in a BST takes $O(\log_2 n)$ in average.
- A BST can be structured so as to have an average access time of $O(n)$. *Can you think of one such BST?*
- Q: Is there a way to *guarantee a worst-case access time of $O(\log_2 n)$ per node* or can we find a way to *guarantee a BST depth of $O(\log_2 n)$* ?
- A: ***AVL Trees***

Definition

An *AVL tree* is a *BST* with the following *balance condition*:

for each node in the BST, the height of left and right sub-trees can differ by at most 1, or

$$\left| h_{N_L} - h_{N_R} \right| \leq 1.$$

Remarks on Balance Condition

- *Balance condition must be easy to maintain:*
 - This is the reason, for example, for the balance condition's not being as follows: the height of left and right sub-trees of each node have the same height.
- *It ensures the depth of the BST is $O(\log_2 n)$.*
- The *height information is stored* as an additional field in `BTNodeType`.

Structure of an AVL Tree

```
struct BTNodeType {  
    infoType *data;  
    unsigned int height;  
    struct BTNodeType *left;  
    struct BTNodeType *right;  
}
```

Rotations


Definition:

- *Rotation* is the operation performed on a BST to restore its AVL property lost as a result of an insert operation.
- We consider the node α whose new balance violates the AVL condition.


Rotation

- Violation of AVL condition
- The AVL condition violation may occur in four cases:
 - Insertion into *left subtree of the left child* (L/L)
 - Insertion into *right subtree of the left child* (R/L)
 - Insertion into *left subtree of the right child* (L/R)
 - Insertion into *right subtree of the right child* (R/R)
- The outside cases 1 and 4 (i.e., L/L and R/R) are fixed by a *single rotation*.
- The other cases (i.e., R/L and L/R) need two rotations called *double rotation* to get fixed.
- These are fundamental operations in balanced-tree algorithms.

Single Rotation (L/L)




before single rotation



after single rotation

$\alpha \equiv k2$
node


Single Rotation (R/R)



before single rotation

$\alpha \equiv k1$



node



after single rotation

Double Rotation (R/L)


$\alpha \equiv k1$
node



Single rotation cannot fix the AVL condition violation!!!

Double Rotation (R/L)

$\alpha \equiv k1$
node



The symmetric case (L/R) is handled similarly left as an exercise to you!


Constructing an AVL Tree – Animation

48

48


Constructing an AVL Tree – Animation

48 16




Constructing an AVL Tree – Animation

48 16 24




Constructing an AVL Tree – Animation

48 16 24




Constructing an AVL Tree – Animation

48 16 24 20




Constructing an AVL Tree – Animation

48 16 24 20 8




Constructing an AVL Tree – Animation

48 16 24 20 8 12




Constructing an AVL Tree – Animation

48 16 24 20 8 12




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18 96




Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18 96



Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18 96 64 17 60 98 68 84 36 30



Height versus Number of Nodes

- The *minimum number* of nodes in an AVL tree recursively relates to the height of the tree as follows:

$$S(h) = S(h-1) + S(h-2) + 1;$$

$$\text{Initial Values: } S(0)=1; S(1)=2$$

Homework: Solve for $S(h)$ as a function of h !

Splay Trees

Motivation for Splay Trees

- We are looking for a data structure where, *even though some worst case ($O(n)$) accesses may be possible, m consecutive tree operations starting from an empty tree (inserts, finds and/or removals) take $O(m \cdot \log_2 n)$.*
- Here, the main idea is to assume that, *$O(n)$ accesses are not bad as long as they occur relatively infrequently.*
- Hence, we are looking for *modifications of a BST per tree operation that attempts to minimize $O(n)$ accesses.*


Splaying

- The underlying idea of splaying is to *move a deep node accessed upwards to the root*, assuming that it will be accessed in the near future again.
- While doing this, other deep nodes are also carried up to smaller depth levels, making the average depth of nodes closer to $O(\log_2 n)$.

Splaying

- Splaying is similar to bottom-up AVL rotations
- If a node *X is the child of the root R*,
 - then we *rotate only X and R, and this is the last rotation performed*.else consider *X*, its *parent P* and *grandparent G*.
Two cases and their symmetries to consider
Zig-zag case, and
Zig-zig case.



Zig-zag case




This is the same operation as an AVL double rotation in an R/L violation.

Zig-zig case


LC(P): left child of node P
RC(P): right child of node P




Animated Example




Animated Example




Animated Example



Animated Example




Animated Example



Node with 6 accessed!

Animated Example



Node with 6 accessed!

B-Trees

Motivation for B-Trees

- Two technologies for providing memory capacity in a computer system
 - *Primary (main) memory (silicon chips)*
 - *Secondary storage (magnetic disks)*
 - Primary memory
 - 5 orders of magnitude (i.e., about 10^5 times) *faster*,
 - 2 orders of magnitude (about 100 times) *more expensive*, and
 - by at least 2 orders of magnitude *less in size*
- than secondary storage due to mechanical operations involved in magnetic disks.

Motivation for B-Trees

- During one disk read or disk write (4-8.5msec for 7200 RPM disks), MM can be accessed about 10^5 times (100 nanosec per access).
- To reimburse (compensate) for this time, at each disks access, *not a single item*, but one or more *equal-sized pages* of items (each page 2^{11} - 2^{14} bytes) are accessed.
- We need some data structure to store these *equal sized pages* in MM.
- *B-Trees*, with their *equal-sized leaves (as big as a page)*, are suitable data structures for storing and performing regular operations on paged data.

B-Trees

- A *B-tree* is a rooted tree with the following properties:
- Every node x has the following fields:
 - $n[x]$, the number of keys currently stored in x .
 - the $n[x]$ keys themselves, in *non-decreasing order*, so that

$$key_1[x] \leq key_2[x] \leq \dots \leq key_{n[x]}[x] ,$$

- $leaf[x]$, a boolean value, true if x is a leaf.

B-Trees

- *Each internal node has $n[x]+1$ pointers, $c_1[x], \dots, c_{n[x]+1}[x]$, to its children. Leaf nodes have no children, hence no pointers!*
- The keys separate the ranges of keys stored in each subtree: if k_i is any key stored in the subtree with root $c_i[x]$, then
$$k_1 \leq key_1[x] \leq k_2 \leq key_2[x] \leq \dots \leq key_{n[x]+1}[x] \leq k_{n[x]+1} .$$
- *All leaves have the same depth, h , equal to the tree's height.*

B-Trees

- There are lower and upper bounds on the number of keys a node may contain. These bounds can be expressed in terms of a fixed integer $t \geq 2$ called the *minimum degree* of the B-Tree.
 - Lower limits
 - All *nodes but the root* has *at least $t-1$* keys.
 - Every *internal node but the root* has *at least t children*.
 - A non-empty tree's **root** must have *at least one key*.

B-Trees

- Upper limits
 - Every *node* can contain *at most $2t-1$ keys*.
 - Every *internal node* can have *at most $2t$ children*.
 - A node is defined to be full if it has exactly *$2t-1$ keys*.
- For a *B-tree* of minimum degree $t \geq 2$ and n nodes

$$h \leq \log_t \frac{n+1}{2}$$

Basic Operations on B-Trees

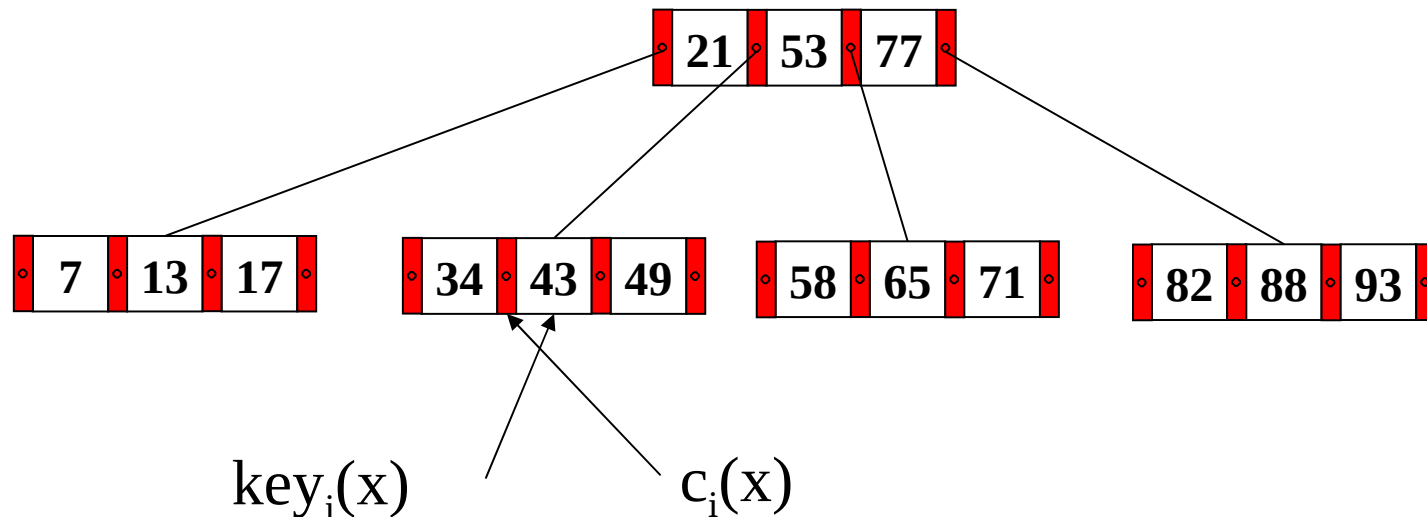
- *B-tree search*
- *B-tree insert*
- *B-tree removal*

Disk Operations in B-Tree operations

- Suppose x is a pointer to an object.
- It is accessible if it is in the main memory.
- If it is on the disk, it needs to be transferred to the main memory to be accessible. This is done by *DISK_READ(x)*.
- To save any changes made to any field(s) of the object pointed to by x , a *DISK_WRITE(x)* operation is performed.

Search in B-Trees

- Similar to search in BSTs with the exception that instead of a binary, a multi-way ($n[x]+1$ -way) decision is made.



Search in B-Trees

B-tree-Search(x,k)

```
{ i=1;
  while ( $i \leq n[x]$  and  $k > key_i[x]$ )  $i++$ ;
  if ( $i \leq n[x]$  and  $k = key_i[x]$ ) // if key found
    return (x,i);
  if (leaf[x]) // if key not found at a leaf
    return NULL;
  else {DISK_READ( $c_i[x]$ ); // if key < keyi[x]
        return B-tree-Search( $c_i[x]$ ,k);}
}
```

Insertion in B-Trees


- Insertion into a B-tree is more complicated than that into a BST, since the creation of a new node to place the new key may violate the B-tree property of the tree.
- Instead, the key is put *into a leaf node x if it is not full*.
- If full, a *split* is applied, which splits a full node (with $2t-1$ keys) at its *median key*, $key_t[x]$, into two nodes with $t-1$ keys each.
- $key_t[x]$ moves up into the parent of x and identifies the split point of the two new trees.

Insertion in B-Trees

- A *single-pass insertion* starts at the root traversing *down to the leaf* into which the key is to be inserted.
- On the path down, *all full nodes are split* including a full leaf that also guarantees a parent with an available position for the median key of a full node to be placed.


Insertion in B-Trees: Example

69 inserted...




Insertion in B-Trees: Example

69 inserted...




Insertion in B-Trees: Example

69 inserted...



Insertion in B-Trees: Example

69 inserted...



Insertion in B-Trees: B-tree-Insert


```
B-tree-Insert(T,k)
{
  r=root[T];
  if (n[r] == 2t-1) {
    s=malloc(new-B-tree-node);
    root[T]=s;
    leaf[s]=false;
    n[s]=0;
    c1[s]=r;
    B-tree-Split-Child(s,1,r);
    B-tree-Insert-Nonfull(s,k); }
  else      B-tree-Insert-Nonfull(r,k);
}
```

Insertion in B-Trees: B-tree-Split-Child


B-tree-Split-Child(x,i,y)

```
{  z=malloc(new-B-tree-node);
   leaf[z]=leaf[y];
   n[z]=t-1;
   for (j = 1; j < t) key[z]=keyj+t[y];      A
   if (!leaf[y])
   for (j = 1; j <= t; j++) cj[z]=cj+t[y];    B      }
   n[y]=t-1;
   for (j=n[x]+1; j>=i+1; j--) cj+1[x]=cj[x];  C      }
   ci+1[x]=z;                                D
   for (j=n[x]; j>=i; j--) keyj+1[x]=keyj[x];  E
   keyi[x]=keyi[y]; n[x]++;                    F
   DISK_WRITE(y);
   DISK_WRITE(z);
   DISK_WRITE(x);
}
```


B-tree-Split-Child: Example



B-tree-Split-Child: Example



B-tree-Split-Child: Example



Insertion in B-Trees: B-tree-Insert-Nonfull

B-tree-Insert-Nonfull(x,k)

```
{  i=n[x];
   if (leaf[x])
       while (i≥1 and k < keyi[x]) {keyi+1[x]=keyi[x]; ;
       i--;}
   keyi+1[x]=k;
   n[x]++;
   DISK_WRITE(x);
else {
   while (i≥1 and k < keyi[x]) i--;
   i++;
   DISK_READ(ci[x]);
   if (n[ci[x]]==2t-1) {
       B-tree-Split-Child(x,i, ci[x]);
       if (k > keyi[x]) i++;
   }
   B-tree-Insert-Nonfull(ci[x],k);
}
```

if x is a leaf
then place key in x;
write x on disk;
else find the node (root of
subtree) key goes to;
read node from disk;
if node full
split node at key's
position;
recursive call with
node split and key;

Removing a key from a B-Tree

- Removal in B-trees is different than insertion only in that *a key may be removed from any node, not just from a leaf*.
- As the insertion algorithm splits any full node down the path to the leaf to which the key is to be inserted, a recursive removal algorithm may be written to ensure that for any call to removal on a node x , the number of keys in x is at least the minimum degree t .

Various Cases of Removing a key from a B-Tree

1. If the key k is in node x and x is a leaf, remove the key k from x .
2. If the key k is in node x and x is an internal node, then
 - a. If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' , and replace k by k' in x . Finding k' and deleting it can be performed in a single downward pass.

Various Cases of Removal a key from a B-Tree

- a. Symmetrically, if the child z that follows k in node x has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k' , and replace k by k' in x . Finding k' and deleting it can be performed in a single downward pass.
- b. Otherwise, if both y and z have only $t-1$ keys, merge k and all of z into y so that x loses both k and the pointer to z and y now contains $2t-1$ keys. Free z and recursively delete k from y .

Various Cases of Removal a key from a B-Tree

1. If k is not present in internal node x , determine root $c_i[x]$ of the subtree that must contain k , if k exists in the tree. If $c_i[x]$ has only $t-1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then finish by recursing on the appropriate child of x .

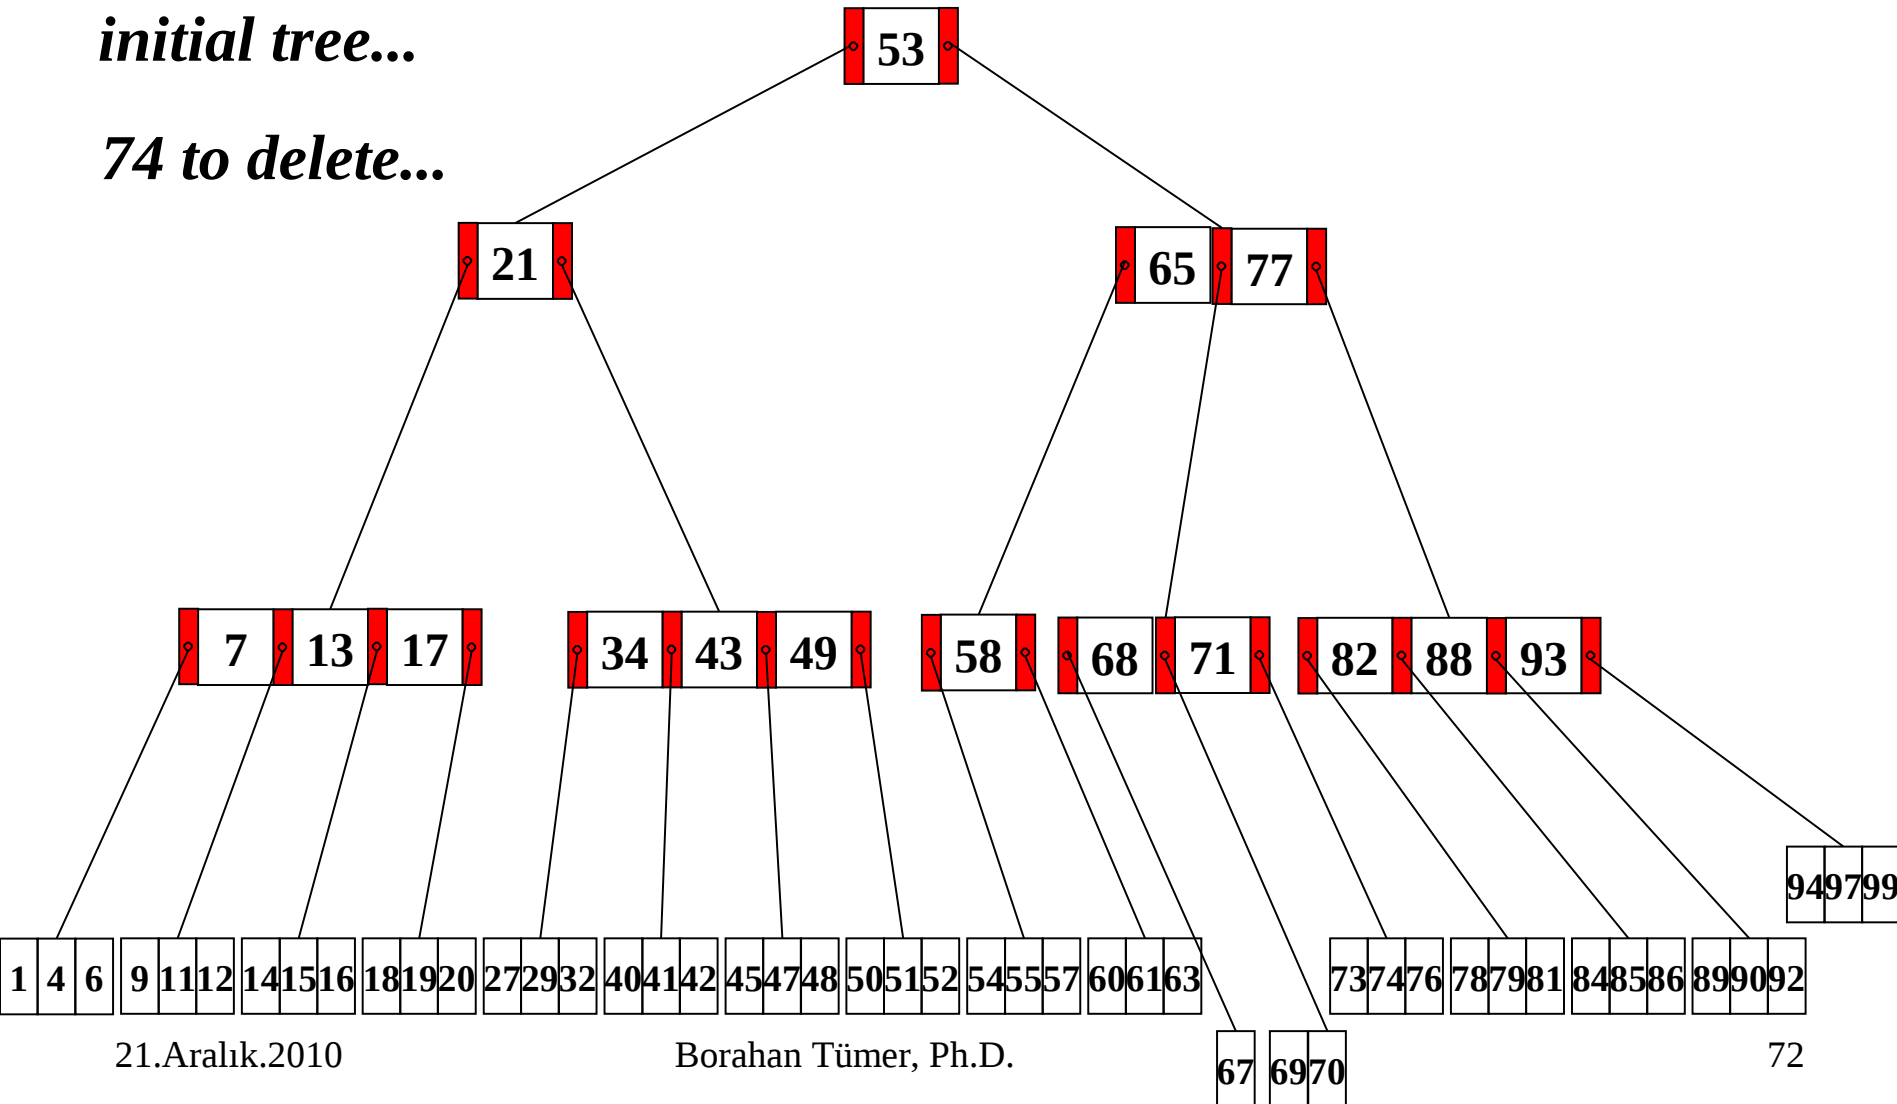
Various Cases of Removal a key from a B-Tree

- a. If $c_i[x]$ has only $t-1$ keys but has an immediate sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x down into $c_i[x]$, moving a key from $c_i[x]$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $c_i[x]$.
- b. If $c_i[x]$ and both of $c_i[x]$'s immediate siblings have $t-1$ keys, merge $c_i[x]$ with one sibling, which involves moving a key from x down into the new merged node to become the *median key* for that node.

Removal in B-Trees: Example

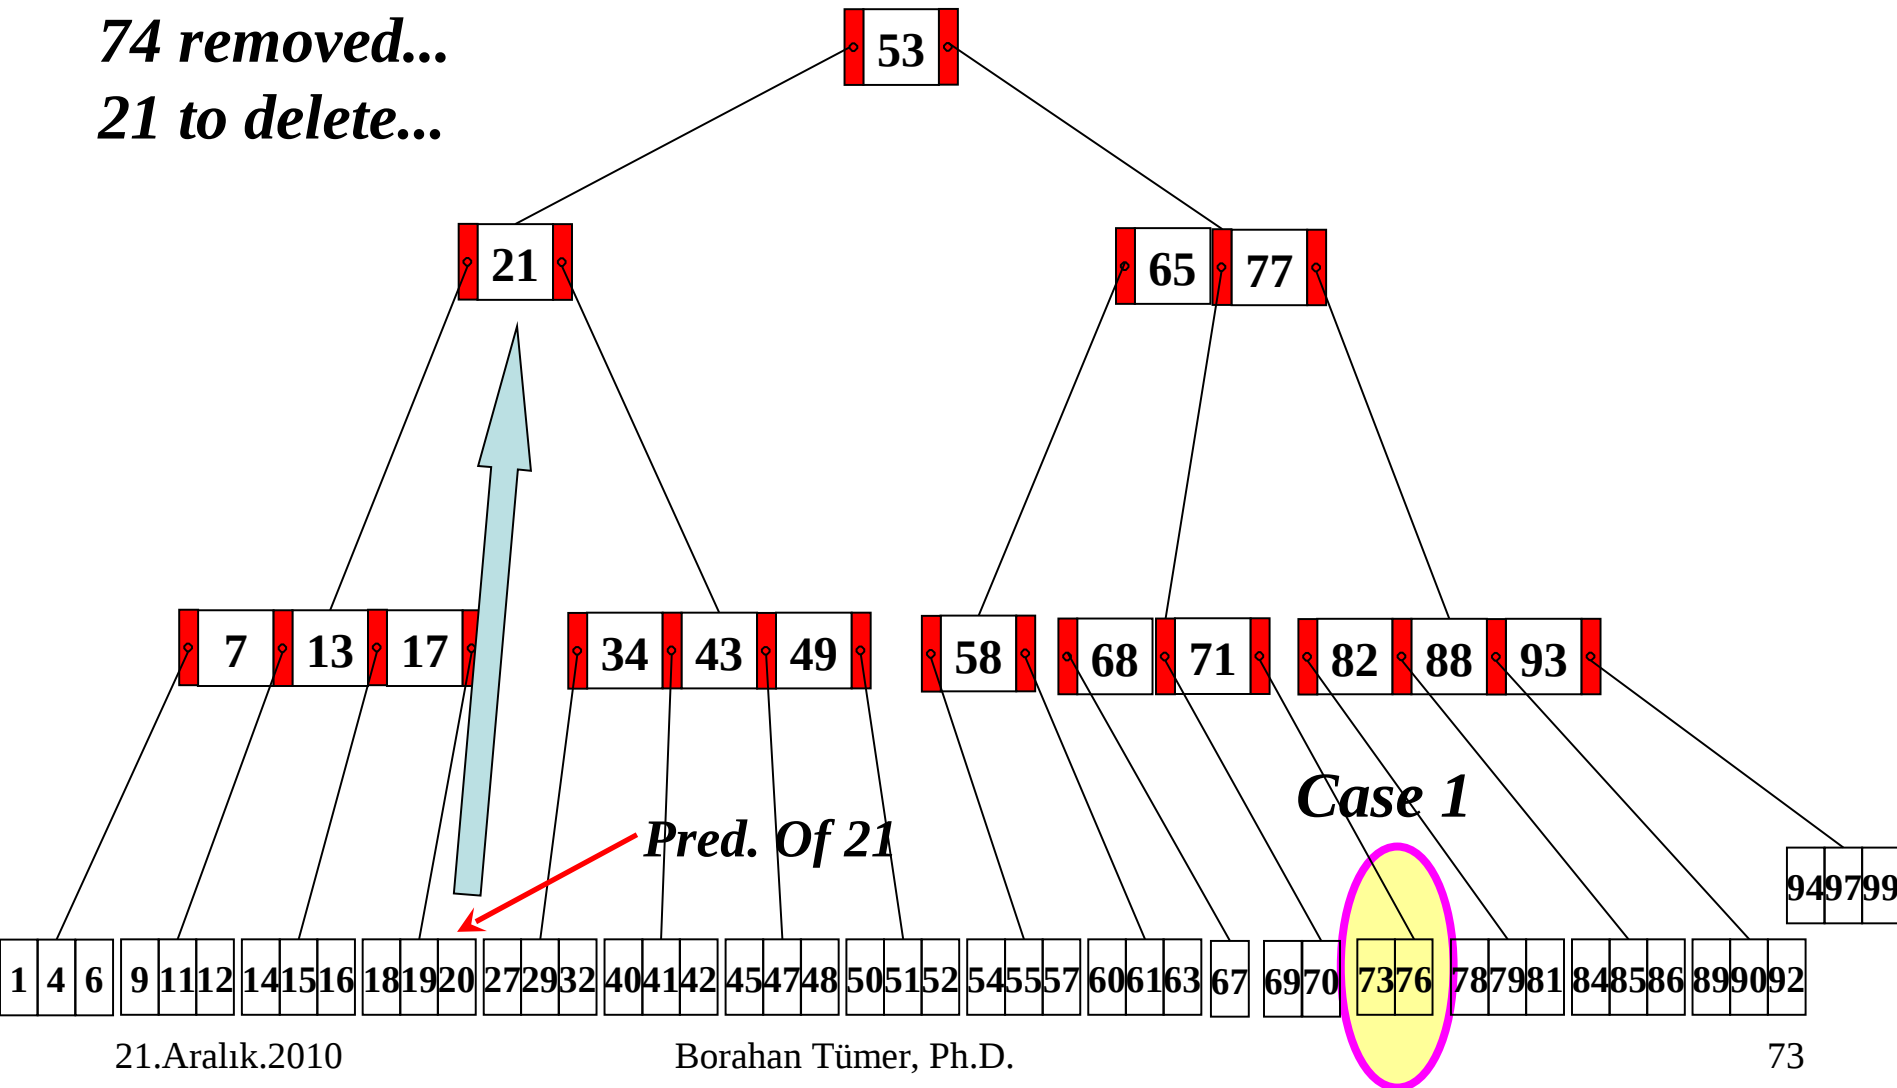
initial tree...

74 to delete...



Removal in B-Trees: Example


74 removed...
21 to delete...



Removal in B-Trees: Example

21 removed...


65 to remove...



Removal in B-Trees: Example

*65 removed...
69 to remove...*


*Case 2b recursively
followed by case 3a*



Removal in B-Trees: Example

69
removed...

*Case 2c recursively
followed by case 1*



Red-Black Trees

- A Red-Black tree (RBT) is a BST with an additional bit “color” information in each node.
- Coloring of nodes is based upon some rules defined in the next slide.
- By putting limits to the coloring of nodes on a path from root to a leaf, RBTs control the length of any such paths and provide an approximate balance. This balance property of RBTs is their desired property.